

Linear Regression with Noise and Regularization

Created by Philip Graff for EN.601.475/675 Spring 2020

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression, Lasso, Ridge
```


Setup

Set the dataset size as well as true model and noise parameters.

```
In [2]: N = 20
x = np.random.rand(N,1)
a = 0.2
b = 0.4
sigma = 0.1
y = a + b * x + np.random.randn(N,1)*sigma
```

Plot the Data

```
In [3]: plt.plot(x, y, '.b')
plt.xlim([0,1])
plt.ylim([0,1])
plt.xlabel('Input')
plt.ylabel('Output')
plt.grid()
plt.show()
```




Fit a Model

Fit a regression model without any regularization and plot the fit versus the data and true model.

```
In [4]: func = LinearRegression()
func.fit(x, y)
xp = np.linspace(0,1,10).reshape(-1,1)
yt = a + b * xp
yp = func.predict(xp)
```

```
In [5]: plt.figure(figsize=(7,6))
plt.plot(x, y, '.b', label='Samples')
plt.plot(xp, yt, '--b', alpha=0.5, label='True')
plt.plot(xp, yp, '-r', alpha=0.75, label='Fit')
plt.xlim([0,1])
plt.ylim([0,1])
plt.xlabel('Input')
plt.ylabel('Output')
plt.grid()
plt.legend(loc='upper left')
plt.show()
```




How Stable is this w.r.t. Noise?

Create multiple datasets, each with a different noise realization and see how the fit varies over the different examples of noise.

```
In [6]: def get_sample_fit(n, sigma, a, b, reg=None, alpha=1.0):
        x = np.random.rand(n,1)
        y = a + b * x + np.random.randn(n,1)*sigma
        if reg is None or reg == 0:
            m = LinearRegression()
            m.fit(x, y)
            return m
        elif reg == 1:
            m = Lasso(alpha=alpha)
            m.fit(x, y)
            return m
        elif reg == 2:
            m = Ridge(alpha=alpha)
            m.fit(x, y)
            return m
        else:
            raise ValueError()
```

```
In [7]: fig, ax = plt.subplots(1,1,figsize=(8,6))
        for _ in range(100):
            func = get_sample_fit(N, sigma, a, b)
            yp = func.predict(xp)
            ax.plot(xp, yp, '-r', alpha=0.1)
        ax.plot(xp, yt, '-b', lw=2)
        ax.set_xlim([0,1])
        ax.set_ylim([0,1])
        ax.set_xlabel('Input')
        ax.set_ylabel('Output')
        ax.grid()
        plt.show()
```



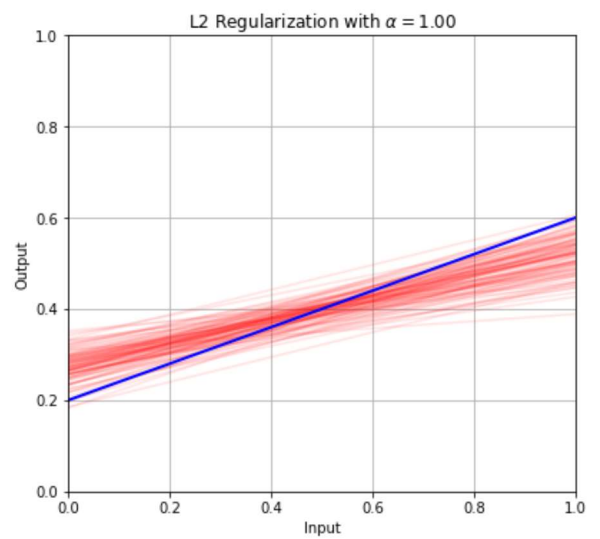

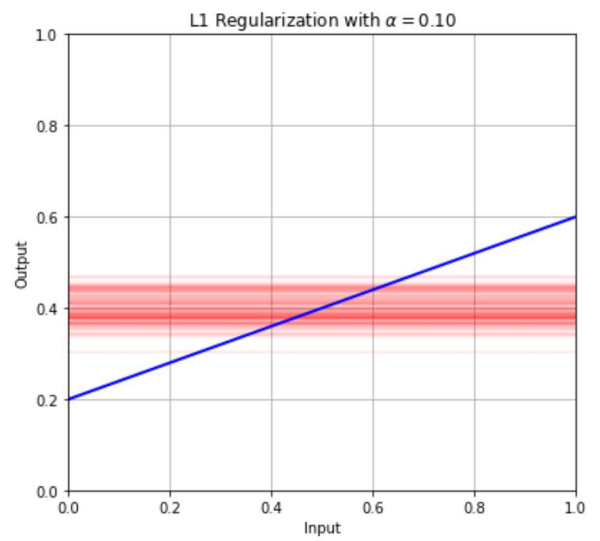
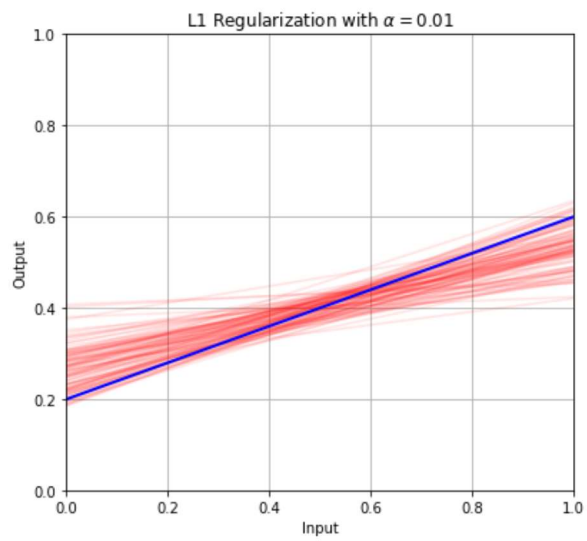

Add Regularization

Repeat, but now with L1 and L2 regularization

```

In [8]: alpha_set = [[0.01, 0.1],[0.1,1.0]]
fig, ax = plt.subplots(3,2,figsize=(14,20))
for _ in range(100):
    func = get_sample_fit(N, sigma, a, b, reg=0)
    yp = func.predict(xp)
    ax[0,0].plot(xp, yp, '-r', alpha=0.1)
ax[0,0].plot(xp, yt, '-b', lw=2)
ax[0,0].set_xlim([0,1])
ax[0,0].set_ylim([0,1])
ax[0,0].set_xlabel('Input')
ax[0,0].set_ylabel('Output')
ax[0,0].set_title('No Regularization ')
ax[0,0].grid()
ax[0,1].axis('off')
for i in range(1,3):
    for j in range(2):
        alpha = alpha_set[i-1][j]
        for _ in range(100):
            func = get_sample_fit(N, sigma, a, b, reg=i, alpha=alpha)
            yp = func.predict(xp)
            ax[i,j].plot(xp, yp, '-r', alpha=0.1)
ax[i,j].plot(xp, yt, '-b', lw=2)
ax[i,j].set_xlim([0,1])
ax[i,j].set_ylim([0,1])
ax[i,j].set_xlabel('Input')
ax[i,j].set_ylabel('Output')
ax[i,j].set_title(r'L%d Regularization with $\alpha$=%.2f'%(i,alpha))
ax[i,j].grid()
plt.show()

```



In []: